



DNN - MNIST for handwritten character recognition

<https://youtu.be/pnSBZ6TEVjY>

- Automated handwritten digit recognition is widely used today - from recognizing zip codes on mail envelopes to recognizing amounts written on bank checks
- We want to use a *MNIST* dataset in TensorFlow
- MNIST is a database of 55000 training examples (`mnist.train`) of labeled handwritten digits '0' to '9' from 250 writers, formatted in 28x28 px matrix (each pixel is encoded as a unsigned byte). The dataset includes also 10000 labeled testing samples (`mnist.test`) from a different set of writers
- If 'x' is the image and 'y' is the label, then `mnist.train.x` and `mnist.train.y` are two tensors



rank is 2, shape is [55000, 784]



rank is 2, shape is [55000, 10]

- If our model does not take into consideration the 2D structure of the images, it is convenient to flatten a 28x28 square matrix as a 784 pixel vector. The pixel order is not important because all images are consistent
- Let's create the computation graph, or the DNN model

```

deep-net.py
1  # Python
2  import os
3  os.environ ['TF_CPP_MIN_LOG_LEVEL']='2'
4
5  import tensorflow as tf
6  from tensorflow.examples.tutorials.mnist import input_data
7  mnist = input_data.read_data_sets("/tmp/data/", one_hot=True)
8
9  n_nodes_hl1 = 500
10 n_nodes_hl2 = 500
11 n_nodes_hl3 = 500
12
13 n_classes = 10 #the number of output nodes
14 batch_size = 100
15
16 x = tf.placeholder('float', [None, 784]) #these are the 28x28 px samples
17 y = tf.placeholder('float') #these are the labels
18
19 # this function describes the feedforward process
20 def neural_network_model(data):
21     # out = (weights * in) + bias
22     hidden_1_layer = {'weights': tf.Variable(tf.random_normal([784, n_nodes_hl1])),
23                       'biases': tf.Variable(tf.random_normal([n_nodes_hl1]))}
24
25     hidden_2_layer = {'weights': tf.Variable(tf.random_normal([n_nodes_hl1, n_nodes_hl2])),
26                       'biases': tf.Variable(tf.random_normal([n_nodes_hl2]))}
27
28     hidden_3_layer = {'weights': tf.Variable(tf.random_normal([n_nodes_hl2, n_nodes_hl3])),
29                       'biases': tf.Variable(tf.random_normal([n_nodes_hl3]))}
30
31     output_layer = {'weights': tf.Variable(tf.random_normal([n_nodes_hl3, n_classes])),
32                     'biases': tf.Variable(tf.random_normal([n_classes]))}
33
34     l1 = tf.add(tf.matmul(data, hidden_1_layer['weights']), hidden_1_layer['biases'])
35     l1 = tf.nn.relu(l1)
36     l2 = tf.add(tf.matmul(l1, hidden_2_layer['weights']), hidden_2_layer['biases'])
37     l2 = tf.nn.relu(l2)
38     l3 = tf.add(tf.matmul(l2, hidden_3_layer['weights']), hidden_3_layer['biases'])
39     l3 = tf.nn.relu(l3)
40     output = tf.add(tf.matmul(l3, output_layer['weights']), output_layer['biases'])
41     return (output)

```

- Then add the training function

```

44 def train_neural_network(x):
45     prediction = neural_network_model(x)
46     cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=prediction, labels=y))
47     optimizer = tf.train.AdamOptimizer().minimize(cost) #learning rate is 0.001
48
49     hm_epochs = 10 # cycles feedforward + backpropagation
50
51     with tf.Session() as sess:
52         sess.run(tf.global_variables_initializer())
53
54         # network training:
55         for epoch in range(hm_epochs):
56             epoch_loss = 0
57             for _ in range(int(mnist.train.num_examples/batch_size)):
58                 epoch_x, epoch_y = mnist.train.next_batch(batch_size)
59                 _, c = sess.run([optimizer, cost], feed_dict = {x: epoch_x, y: epoch_y})
60                 epoch_loss += c
61             print('Epoch', epoch, 'completed oput of', hm_epochs, 'loss:', epoch_loss)
62
63         # print results
64         correct = tf.equal(tf.argmax(prediction,1), tf.argmax(y,1))
65         accuracy = tf.reduce_mean(tf.cast(correct, 'float'))
66         print('Accuracy:', accuracy.eval({x: mnist.test.images, y:mnist.test.labels}))
67
68     train_neural_network(x)

```

- Run the program with Python3

```
Extracting /tmp/data/train-images-idx3-ubyte.gz
Extracting /tmp/data/train-labels-idx1-ubyte.gz
Extracting /tmp/data/t10k-images-idx3-ubyte.gz
Extracting /tmp/data/t10k-labels-idx1-ubyte.gz
Epoch 0 completed oput of 10 loss: 1821681.66306
Epoch 1 completed oput of 10 loss: 431706.899899
Epoch 2 completed oput of 10 loss: 238944.139815
Epoch 3 completed oput of 10 loss: 141129.630415
Epoch 4 completed oput of 10 loss: 87271.4430515
Epoch 5 completed oput of 10 loss: 57762.353711
Epoch 6 completed oput of 10 loss: 38104.9504944
Epoch 7 completed oput of 10 loss: 28384.9141481
Epoch 8 completed oput of 10 loss: 22437.0027955
Epoch 9 completed oput of 10 loss: 19794.682209
Accuracy: 0.9475
[Finished in 137.7s]
```

The accuracy achieved is 0.9475%.

Notes

- `tf.placeholder` represents a symbolic variable for the input
- For example, with `x = tf.placeholder(tf.float32, [None, 784])` we can represent a 2D tensor of floating-point numbers with a `[None, 784]` shape
- `tf.Variable` is a modifiable tensor that stores the model parameters
- For example, with `W=tf.Variable(tf.zeros([784, 10]))` we set the initial values to the weights from any image and the output classes